

Homework week 5

2022-06-18

Question 11.1

Question 11.1.1

Answer:

To run a stepwise regression, we need to use the built-in `step()` function from `stats` package.
First, read the data.

```
rm(list = ls())
set.seed(19)
data <- read.table("C:\\Data\\week 5 data-summer\\data 11.1\\uscrime.txt",
                  stringsAsFactors = FALSE, header = TRUE)
head(data)
```

```
##      M So   Ed Po1  Po2    LF   M.F Pop   NW    U1  U2 Wealth Ineq    Prob
## 1 15.1   1  9.1  5.8  5.6 0.510  95.0  33 30.1 0.108 4.1   3940 26.1 0.084602
## 2 14.3   0 11.3 10.3  9.5 0.583 101.2  13 10.2 0.096 3.6   5570 19.4 0.029599
## 3 14.2   1  8.9  4.5  4.4 0.533  96.9  18 21.9 0.094 3.3   3180 25.0 0.083401
## 4 13.6   0 12.1 14.9 14.1 0.577  99.4 157  8.0 0.102 3.9   6730 16.7 0.015801
## 5 14.1   0 12.1 10.9 10.1 0.591  98.5  18  3.0 0.091 2.0   5780 17.4 0.041399
## 6 12.1   0 11.0 11.8 11.5 0.547  96.4  25  4.4 0.084 2.9   6890 12.6 0.034201
##      Time Crime
## 1 26.2011    791
## 2 25.2999   1635
## 3 24.3006    578
## 4 29.9012   1969
## 5 21.2998   1234
## 6 20.9995    682
```

Next, we need to define two models: one includes intercept term only, the other includes all predictors.

```
#define intercept-only model
intercept_only <- lm(Crime ~ 1, data=data)

#define model with all predictors
all <- lm(Crime ~ ., data=data)
```

Then we can implement stepwise regression. The `step()` function will use AIC as its variable selection criteria. The info of the final model is displayed below.

```
stepwise <- step(intercept_only, direction='both',
                scope=list(lower = intercept_only, upper = formula(all)),
                trace=0)

#view the steps taken to get the final model
stepwise$anova
```

```
##      Step Df  Deviance Resid. Df Resid. Dev      AIC
## 1      NA      NA      46    6880928 561.0235
## 2 + Po1 -1 3253301.8      45    3627626 532.9352
## 3 + Ineq -1 739818.6      44    2887807 524.2154
## 4 + Ed -1 587049.8      43    2300757 515.5343
## 5 + M -1 239404.6      42    2061353 512.3701
## 6 + Prob -1 258062.5      41    1803290 508.0839
## 7 + U2 -1 192233.4      40    1611057 504.7859
```

```
#View the final model(scaled data)
summary(stepwise)
```

```
##
## Call:
## lm(formula = Crime ~ Po1 + Ineq + Ed + M + Prob + U2, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -470.68  -78.41  -19.68   133.12   556.23
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -5040.50     899.84  -5.602 1.72e-06 ***
## Po1          115.02       13.75   8.363 2.56e-10 ***
## Ineq         67.65       13.94   4.855 1.88e-05 ***
## Ed          196.47       44.75   4.390 8.07e-05 ***
## M           105.02       33.30   3.154 0.00305 **
## Prob       -3801.84    1528.10  -2.488 0.01711 *
## U2           89.37       40.91   2.185 0.03483 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 200.7 on 40 degrees of freedom
## Multiple R-squared:  0.7659, Adjusted R-squared:  0.7307
## F-statistic: 21.81 on 6 and 40 DF, p-value: 3.418e-11
```

Question 11.1.2

Answer:

To execute a LASSO regression, we need to use package `glmnet`. We also need to scale the first 15 columns in the `uscrime` dataset and use them as our predictors.

To perform LASSO regression, parameter “alpha” in the function `glmnet()` should be set as 1. Since `glmnet()` will return more than one models with different lambda values, we would directly use function `cv.glmnet()` to performs 5-fold cross validation and thus identify the lambda value that produces the smallest mean squared error.

The lambda value that generates smallest mean squared error will be used in our final model.

```
library(glmnet)
```

```
## Loading required package: Matrix
```


```
## Loaded glmnet 4.1-4
```

```
#define response variable  
y <- as.matrix(data$Crime)  
  
#define matrix of predictor variables (scaled)  
x <- data.matrix(scale(data[, 1:15]))
```

```
set.seed(19)  
  
#3-fold cross validation  
model <- cv.glmnet(x, y, alpha = 1, nfolds = 5 )  
  
# lambda value that generate smallest MSE  
best_lambda <- model$lambda.min  
best_lambda
```

```
## [1] 12.21181
```

```
plot(model)
```



Next we will plug in $\lambda = 12.21181$ to `glmnet()` to obtain our final model.

Here are the coefficients of the final model (based on scaled data):

```
final_model <- glmnet(x,y, alpha=1, lambda = best_lambda)

#show coefficients of the final model
coef(final_model)
```

```
## 16 x 1 sparse Matrix of class "dgCMatrix"
##                s0
## (Intercept) 905.085106
## M           82.758328
## So          22.929369
## Ed          115.745467
## Po1         310.430186
## Po2         .
## LF          2.352681
## M.F         49.722388
## Pop         .
## NW          3.571258
## U1          -13.474475
## U2          46.340117
## Wealth     .
## Ineq        175.639893
## Prob       -80.369954
## Time       .
```

Question 11.1.3

Answer:

To run the elastic net regression, we are going to use the `train()` function in `caret` package. The `glmnet()` function will be invoked by `train()` function.

We use `train()` to automatically select the best tuning parameters, α and λ . The `train()` function will test a range of possible α and λ values, and then select the best values for λ and α that will minimize the model's mean squared error. A 5-fold cross validation approach will be applied in the tuning process.

The best α and λ values and the final model will be displayed as follows:

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
# Scale the data frame
data_scaled <- as.data.frame(cbind(scale(data[,1:15]), data[,16]))
names(data_scaled)[16] <- "Crime"
head(data_scaled)
```

##	M	So	Ed	Po1	Po2	LF	M.F
## 1	0.9886930	1.3770536	-1.3085099	-0.9085105	-0.8666988	-1.2667456	-1.12060499
## 2	0.3521372	-0.7107373	0.6580587	0.6056737	0.5280852	0.5396568	0.98341752
## 3	0.2725678	1.3770536	-1.4872888	-1.3459415	-1.2958632	-0.6976051	-0.47582390
## 4	-0.2048491	-0.7107373	1.3731746	2.1535064	2.1732150	0.3911854	0.37257228
## 5	0.1929983	-0.7107373	1.3731746	0.8075649	0.7426673	0.7376187	0.06714965
## 6	-1.3983912	-0.7107373	0.3898903	1.1104017	1.2433590	-0.3511718	-0.64550313

##	Pop	NW	U1	U2	Wealth	Ineq
## 1	-0.09500679	1.943738564	0.69510600	0.8313680	-1.3616094	1.6793638
## 2	-0.62033844	0.008483424	0.02950365	0.2393332	0.3276683	0.0000000
## 3	-0.48900552	1.146296747	-0.08143007	-0.1158877	-2.1492481	1.4036474
## 4	3.16204944	-0.205464381	0.36230482	0.5945541	1.5298536	-0.6767585
## 5	-0.48900552	-0.691709391	-0.24783066	-1.6551781	0.5453053	-0.5013026
## 6	-0.30513945	-0.555560788	-0.63609870	-0.5895155	1.6956723	-1.7044289

##	Prob	Time	Crime
## 1	1.6497631	-0.05599367	791
## 2	-0.7693365	-0.18315796	1635
## 3	1.5969416	-0.32416470	578
## 4	-1.3761895	0.46611085	1969
## 5	-0.2503580	-0.74759413	1234
## 6	-0.5669349	-0.78996812	682

```

# Build the model
set.seed(19)
model <- train(Crime ~.,
               data = data_scaled,
               method = "glmnet",
               trControl = trainControl("cv", number = 5),
)
# Different models tried in the tuning process
model

```

```
## glmnet
##
## 47 samples
## 15 predictors
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 38, 38, 37, 39, 36
## Resampling results across tuning parameters:
##
##  alpha  lambda      RMSE      Rsquared  MAE
##  0.10    0.5261908  275.4363  0.5699113  208.5414
##  0.10    5.2619079  270.8485  0.5427620  207.6630
##  0.10   52.6190793  272.3376  0.5155554  217.4572
##  0.55    0.5261908  275.3085  0.5683242  208.2275
##  0.55    5.2619079  268.9321  0.5451441  207.7069
##  0.55   52.6190793  283.7196  0.5028243  224.2369
##  1.00    0.5261908  275.3899  0.5658953  207.9369
##  1.00    5.2619079  269.4731  0.5439143  208.4011
##  1.00   52.6190793  299.4991  0.4914756  234.4098
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 0.55 and lambda = 5.261908.
```

```
#best parameters
model$bestTune
```

```
##  alpha  lambda
## 5  0.55 5.261908
```

```
best_alpha <- model$bestTune$alpha
best_lambda <- model$bestTune$lambda

# display the final model with best alpha and lambda (scaled data)
final_model <- glmnet(x,y, alpha=best_alpha, lambda = best_lambda)
coef(final_model)
```

```
## 16 x 1 sparse Matrix of class "dgCMatrix"
##                s0
## (Intercept) 905.08511
## M           101.81656
## So           17.53785
## Ed           170.49775
## Po1          288.76389
## Po2           .
## LF           .
## M.F           56.03589
## Pop          -16.69157
## NW           18.51147
## U1           -71.97887
## U2           115.30076
## Wealth        54.77997
## Ineq         240.29885
## Prob         -90.46793
## Time           .
```

Question 12.1

Answer:

I am working at a management consulting firm. Recently one of our clients is making advertisements on the internet, and their management what to know which platform (twitter or facebook) is the best channel for online advertising.

I believe this situation is a very good scenario for A/B testing. Our client can broadcast exactly the same advertisements on both two websites. They can determine which website is better channel by viewing the user click rate.

Question 12.2

Answer:

To find a factorial design of this experiment, we simply need to run function FrF2() with following parameters:

nruns = 16 and nfactors = 10

Here are the results:

```
library(FrF2)
```

```
## Loading required package: DoE.base
```

```
## Loading required package: grid
```

```
## Loading required package: conf.design
```

```
## Registered S3 method overwritten by 'partitions':  
##   method          from  
##   print.equivalence lava
```

```
## Registered S3 method overwritten by 'DoE.base':  
##   method          from  
##   factorize.factor conf.design
```

```
##  
## Attaching package: 'DoE.base'
```

```
## The following objects are masked from 'package:stats':  
##  
##   aov, lm
```

```
## The following object is masked from 'package:graphics':  
##  
##   plot.design
```

```
## The following object is masked from 'package:base':  
##  
##   lengths
```

```
FrF2(nruns = 16, nfactors = 10)
```

```
##      A  B  C  D  E  F  G  H  J  K  
## 1  -1 -1  1 -1  1 -1 -1  1  1 -1  
## 2   1  1  1 -1  1  1  1 -1 -1 -1  
## 3  -1  1 -1  1 -1  1 -1 -1 -1  1  
## 4   1  1 -1 -1  1 -1 -1 -1  1  1  
## 5   1 -1 -1  1 -1 -1  1  1  1  1  
## 6  -1  1 -1 -1 -1  1 -1  1  1 -1  
## 7   1  1  1  1  1  1  1  1  1  1  
## 8   1  1 -1  1  1 -1 -1  1 -1 -1  
## 9  -1  1  1 -1 -1 -1  1  1 -1  1  
## 10 -1 -1 -1 -1  1  1  1  1 -1  1  
## 11  1 -1 -1 -1 -1 -1  1 -1 -1 -1  
## 12  1 -1  1  1 -1  1 -1  1 -1 -1  
## 13 -1  1  1  1 -1 -1  1 -1  1 -1  
## 14 -1 -1  1  1  1 -1 -1 -1 -1  1  
## 15  1 -1  1 -1 -1  1 -1 -1  1  1  
## 16 -1 -1 -1  1  1  1  1 -1  1 -1  
## class=design, type= FrF2
```

Question 13.1

Answer:

1. Binomial: in a plant of manufacturing company, there is a defective rate of finished goods coming down from

the assembly line. Assuming the defective rate is constant all the time and independent for each piece of finished goods produced, the number of defective products out of total products produced will follow a binomial distribution.

2. Geometric: In the same example as binomial distribution, the number of non-defective products produced before each one defective product is produced will follow a geometric distribution.

3. Poisson: The number of calls received per hour at a call center will follow a poisson distribution.

4. Exponential: In the same example as poisson distribution, the time (minutes) between each call received will follow a exponential distribution.

5. Weibull: The mileage that a newly produced car can run before any failure will follow a Weibull distribution. Since the parts of the car can wear out as it runs longer, the likelihood of failure may increases as mileage increases. Therefore, we would expect $k > 1$.